


Lecture 9 part 3

Gradient Descent on Manifold

Weiqing Gu

Gradient descent on manifold

Tangent space:

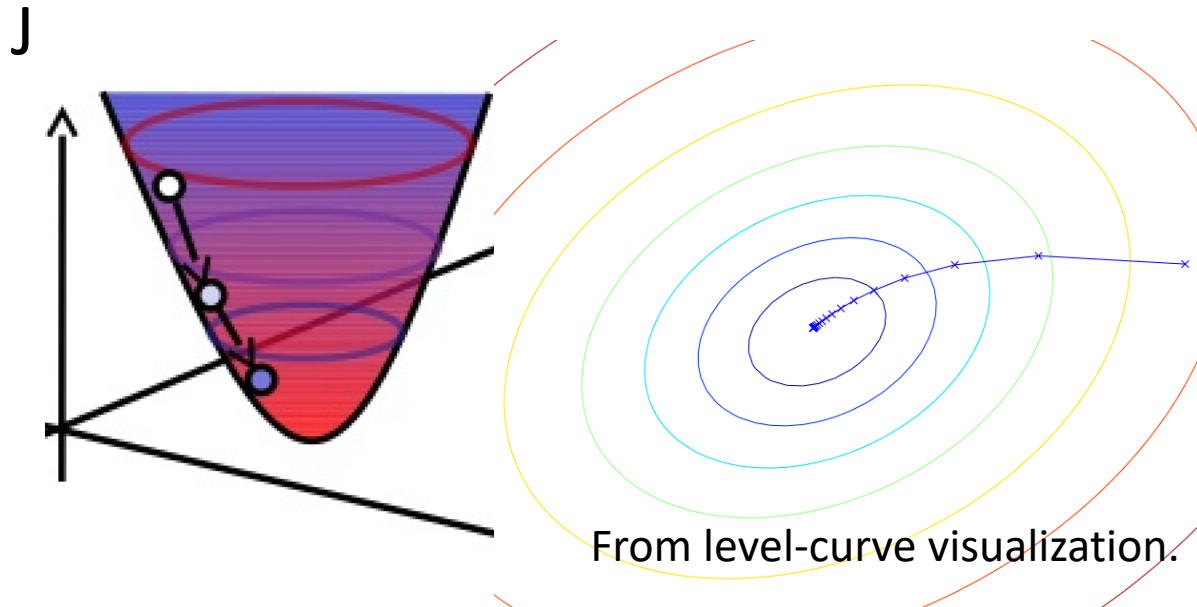


Riemmanian metric: scalar product $\langle u, v \rangle_g$ on the tangent space


First recall gradient descent in \mathbb{R}^n

Use the gradient descent algorithm

- Which starts with some initial θ , and repeatedly performs the update.
- Here α is called the learning rate.
- Geometrically, it repeatedly takes a step in the direction of steepest decrease of J .



Make α smaller if necessary.



Batch Gradient Descent (BGD)


Repeat until convergence {

$$\theta_j := \theta_j + \alpha \underbrace{\sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}}_{-\partial J(\theta)/\partial \theta_j} \quad (\text{for every } j).$$

This is simply gradient descent on the original cost function J.

Remarks:

- 1) **This method looks at every example in the entire training set on every step**, and is called BGD.
- 2) It is well known that gradient descent can be susceptible to local minima in general (see the figure on right), **the optimization problem we have** posed here for linear regression **has only one global**, and no other local, **optima**; thus gradient descent always converges (assuming the learning rate α is not too large) to the global minimum.
- 3) **The key is that our J is a convex quadratic function.**



Stochastic Gradient Descent (SGD)


```
Loop {  
    for i=1 to m, {  
         $\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$       (for every  $j$ ).  
    }  
}
```

Remarks:

- 1) SGD **repeatedly run through the training set, and each time it encounters a training example, it updates the parameters** according to the gradient of the error with respect to that single training example only.
- 2) SGD **may never “converge” to the unique minimum**, and the parameters θ will keep oscillating around the minimum of $J(\theta)$; but **in practice** most of the values near the minimum will be **reasonably good approximations** to the true minimum.

Comparing Batch gradient descent with Stochastic gradient descent

- *For big data*, often the training set is large, *people prefer use stochastic gradient descent* instead of batch gradient descent.
- Since *BGD has to scan thru the entire training set before taking a single step*—a costly operation if m is large—*SGD can start making progress right away*, & continues to make progress with each example it looks at.
- *SGD can run on dynamical data sets*. As data coming, it updates the parameters.
- Often, **SGD gets θ “close” to the minimum much faster than BGD**.
- But SGD gets only approximation solution of θ . This is a **trade off** when dealing with big data.



Now let's see how to extend SGD to a manifold

Consider $f : \mathcal{M} \rightarrow \mathbb{R}$ twice differentiable.

Riemannian gradient: tangent vector at x satisfying

$$\frac{d}{dt}\bigg|_{t=0} f(\exp_x(tv)) = \langle v, \nabla f(x) \rangle_g$$

Riemannian Hessian: based on the Taylor expansion

$$f(\exp_x(tv)) = t\langle v, \nabla f(x) \rangle_g + \frac{1}{2}t^2 v^T [\text{Hess } f(x)] v + O(t^3)$$

Second order Taylor expansion:

$$f(\exp_x(tv)) - f(x) \leq t\langle v, \nabla f(x) \rangle_g + \frac{t^2}{2} \|v\|_g^2 k$$

where k is a bound on the hessian along the geodesic.

Stochastic Gradient descent on manifold

Riemannian approximated gradient: $E_Z(H(z_t, w_t)) = \nabla C(w_t)$
a tangent vector !

Stochastic gradient descent on \mathcal{M} : update

$$w_{t+1} \leftarrow \exp_{w_t}(-\gamma_t H(z_t, w_t))$$

w_{t+1} must remain on \mathcal{M} !

